

# Regression Modelling to Predict the Age of the Abalone based on its Physical Characteristics - Log Transformed Response

Minoli Munasinghe

2023-04-10

Read the dataset from UCI machine learning repository

```
# Read the dataset into a data frame
abalone_df <- read.csv("http://archive.ics.uci.edu/ml/machine-learning-databases/abalone/abalone.data",
# Print first five rows in the data set
head(abalone_df)
```

```
##   V1     V2     V3     V4     V5     V6     V7     V8 V9
## 1 M 0.455 0.365 0.095 0.5140 0.2245 0.1010 0.150 15
## 2 M 0.350 0.265 0.090 0.2255 0.0995 0.0485 0.070  7
## 3 F 0.530 0.420 0.135 0.6770 0.2565 0.1415 0.210  9
## 4 M 0.440 0.365 0.125 0.5160 0.2155 0.1140 0.155 10
## 5 I 0.330 0.255 0.080 0.2050 0.0895 0.0395 0.055  7
## 6 I 0.425 0.300 0.095 0.3515 0.1410 0.0775 0.120  8
```

Add column names to the data set

```
# add column names
names(abalone_df) <- c("sex", "length", "diameter", "height", "weight.whole",
                        "weight.shucked", "weight.viscera", "weight.shell", "rings")
```

Get the number of observations in the data

```
# numbers of rows in the data set
nrow(abalone_df)
```

```
## [1] 4177
```

Get the number of columns in the data

```
ncol(abalone_df)
```

```
## [1] 9
```

The data set has 4177 observations with 9 variables

Get the number of missing values

```
# Check for the missing values  
sum(is.na(abalone_df))
```

```
## [1] 0
```

The data set does not have any missing values

Get the structure of data

```
str(abalone_df)
```

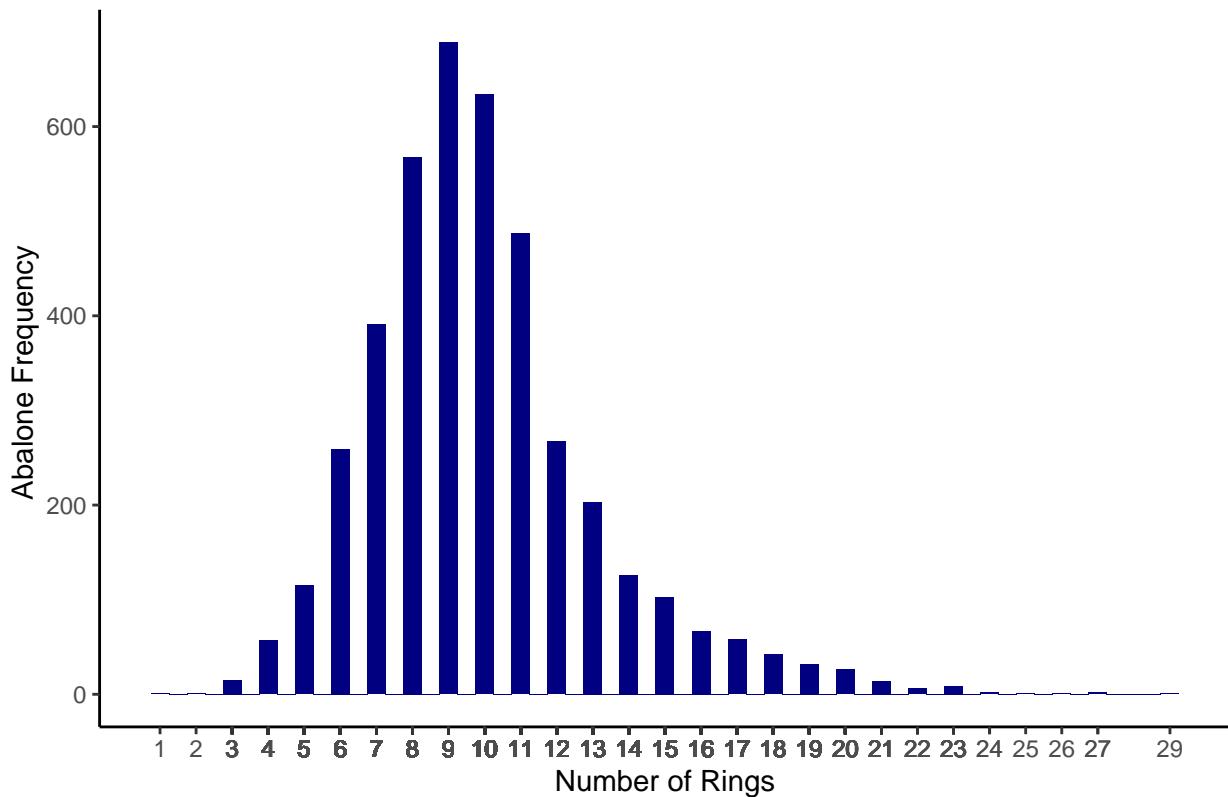
```
## 'data.frame': 4177 obs. of 9 variables:  
## $ sex : chr "M" "M" "F" "M" ...  
## $ length : num 0.455 0.35 0.53 0.44 0.33 0.425 0.53 0.545 0.475 0.55 ...  
## $ diameter : num 0.365 0.265 0.42 0.365 0.255 0.3 0.415 0.425 0.37 0.44 ...  
## $ height : num 0.095 0.09 0.135 0.125 0.08 0.095 0.15 0.125 0.125 0.15 ...  
## $ weight.whole : num 0.514 0.226 0.677 0.516 0.205 ...  
## $ weight.shucked: num 0.2245 0.0995 0.2565 0.2155 0.0895 ...  
## $ weight.viscera: num 0.101 0.0485 0.1415 0.114 0.0395 ...  
## $ weight.shell : num 0.15 0.07 0.21 0.155 0.055 0.12 0.33 0.26 0.165 0.32 ...  
## $ rings : int 15 7 9 10 7 8 20 16 9 19 ...
```

According to the data, there is one categorical variable (sex) and all the other variables are continuous. The response variable of the model is rings.

The distribution of response variable

```
library(ggplot2)  
ggplot(aes(x = rings), data = abalone_df)+  
  geom_histogram(binwidth = 0.5, fill = "navy blue") +  
  scale_x_continuous(breaks = abalone_df$rings)+  
  ylab("Abalone Frequency") +  
  xlab("Number of Rings") +  
  ggtitle("The Frequency Distribution of Abalone Rings") +  
  theme_classic() +  
  theme(plot.title = element_text(hjust = 0.5))
```

## The Frequency Distribution of Abalone Rings



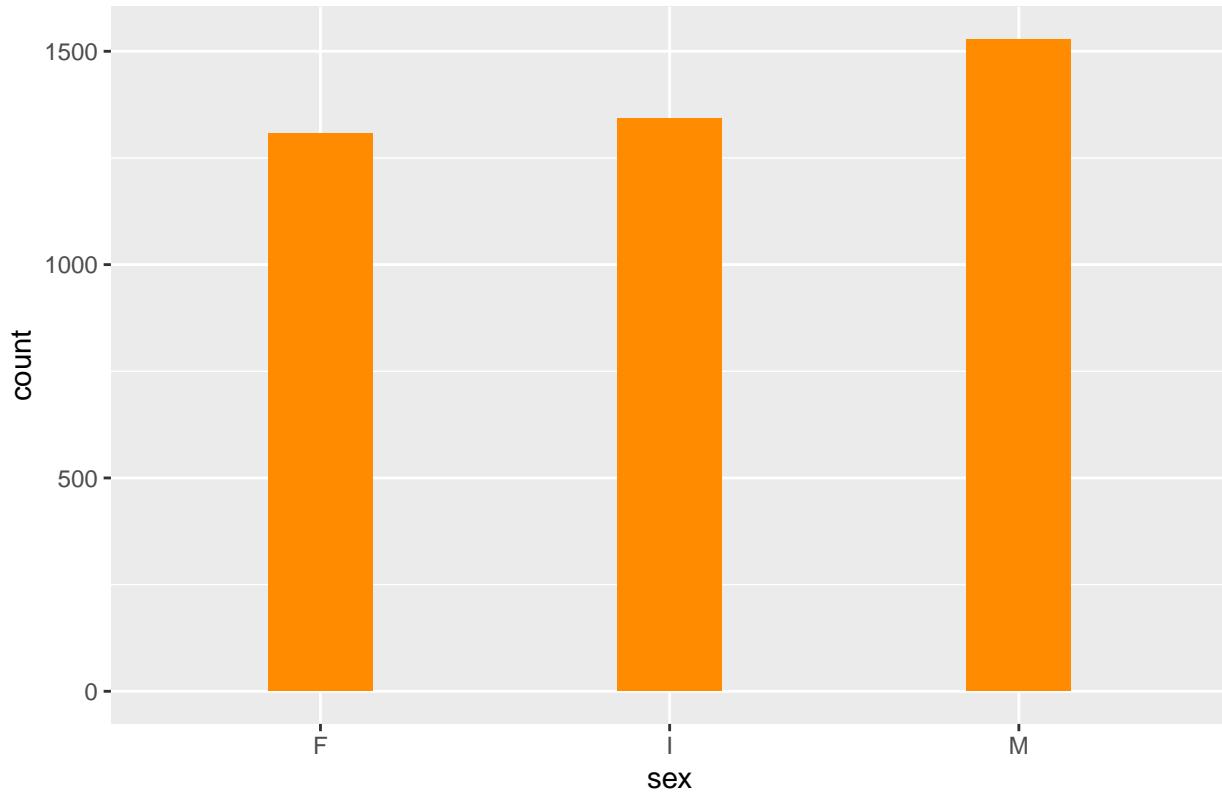
```
library(moments)
skewness(abalone_df$rings)
```

```
## [1] 1.113702
```

The distribution of observations by number of rings is observed using the above histogram and according to its skewness value of 1.1, we can consider that it is moderately positively skewed and the distribution is also not symmetrically distributed around the mean.

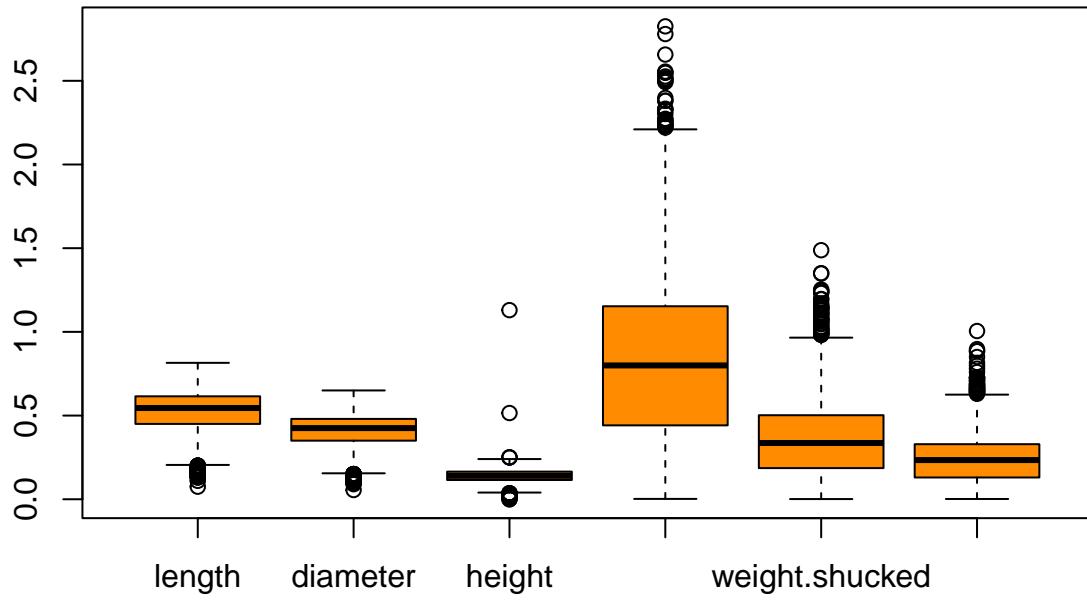
```
library(ggplot2)
ggplot(aes(x = sex), data = abalone_df) +
  geom_bar(stat = "count", width = 0.3, fill = "dark orange") +
  theme(legend.position = "None") +
  ggtitle("Sex Distribution of abalone")
```

## Sex Distribution of abalone



Observe the distribution of other variables

```
boxplot(abalone_df[,c(2,3,4,5,6,8)], col = "dark orange") # boxplot to check possible outliers
```



According to the boxplots drawn for all the other variables, we can see that all variables have outlier values in the data.

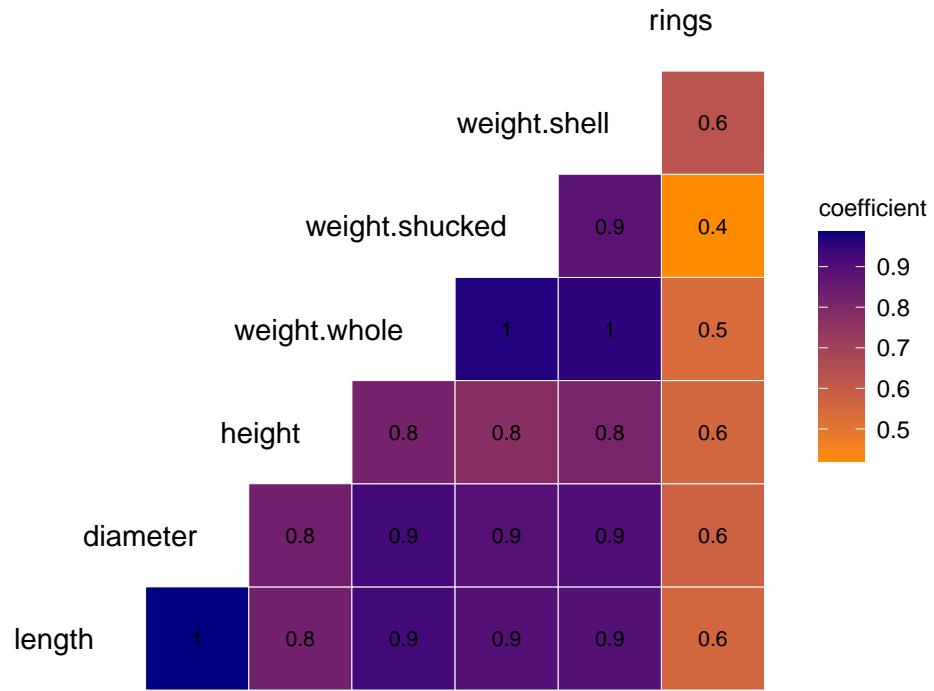
Correlation plot to observe the relationship between variables

```
library(GGally)
```

```
## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg   ggplot2

ggcorr(abalone_df[,c(2,3,4,5,6,8,9)], label = T, label_size = 2.9, hjust = 1, layout.exp = 3) +
  scale_fill_gradient(low = "dark orange", high = "navy blue")

## Scale for fill is already present.
## Adding another scale for fill, which will replace the existing scale.
```

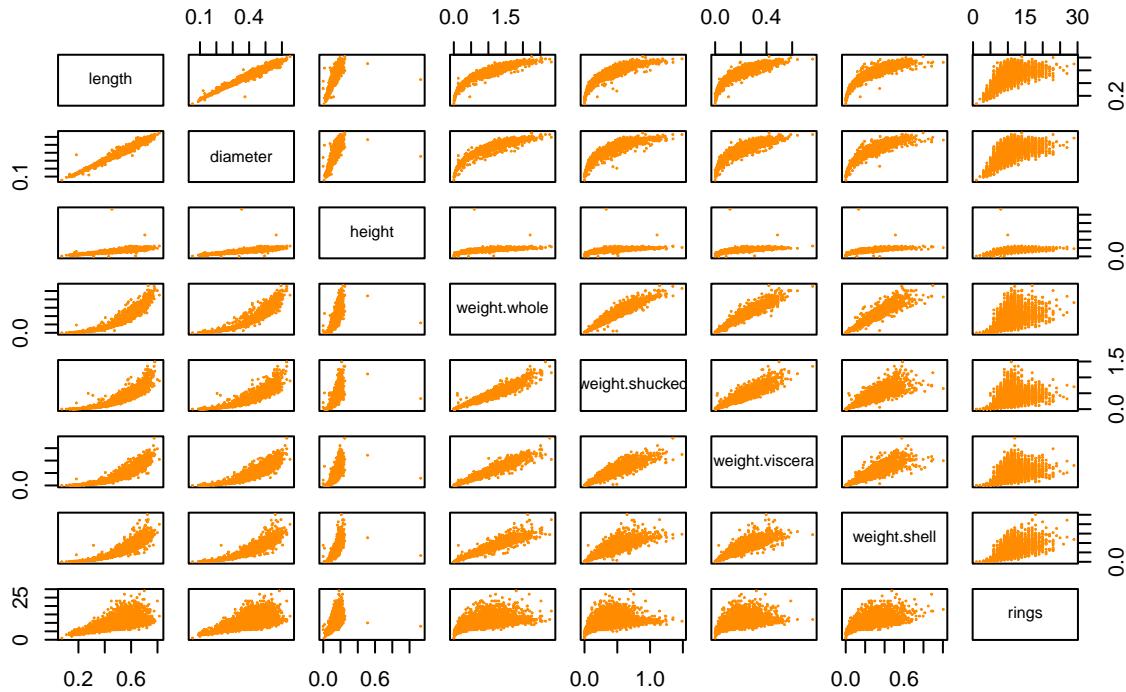


According to the correlation plot, we can see that the predictor variables are highly correlated and therefore multicollinearity exists.

Pairwise Scatterplot

```
# Plot the scatter matrix
pairs(abalone_df[,2:9], main = "Scatterplot Matrix of Abalone Data", col = "dark orange", cex=0.1, cex.main=1.5)
```

## Scatterplot Matrix of Abalone Data



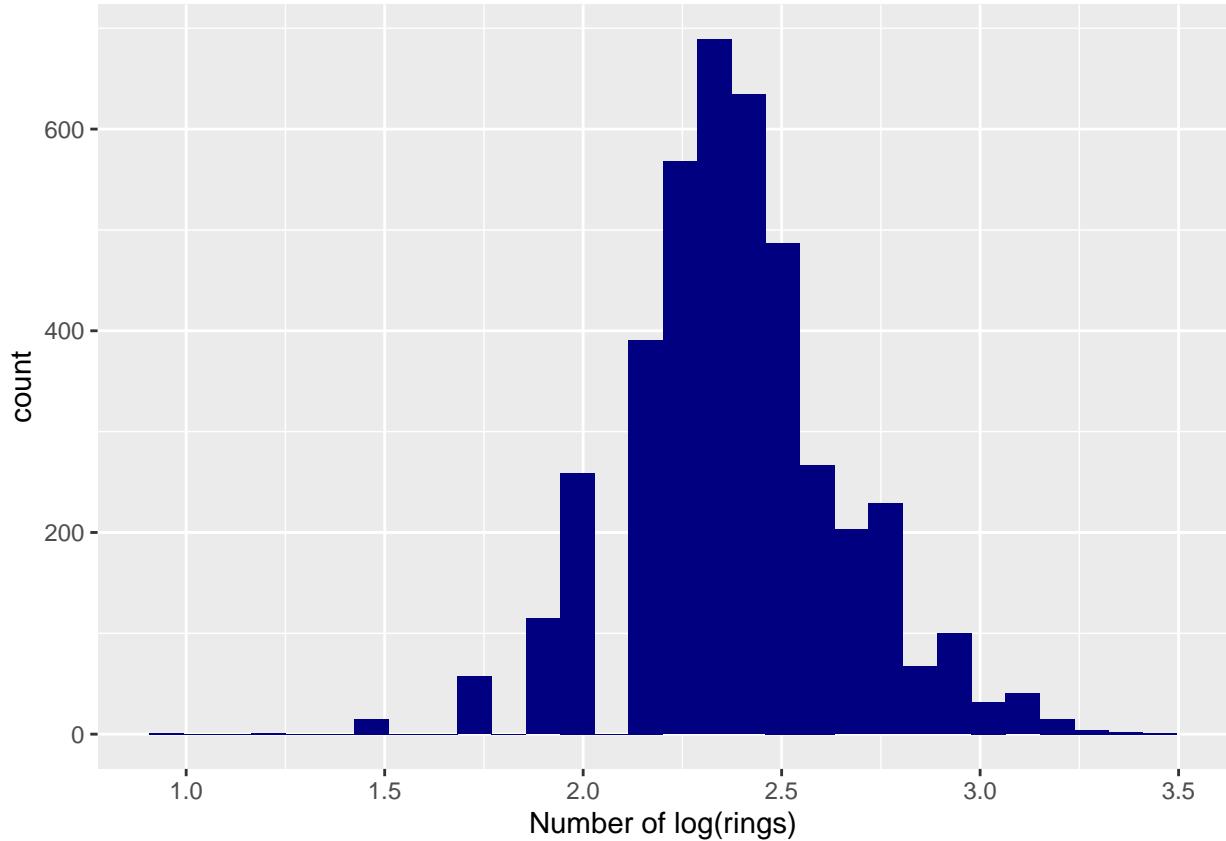
```
abalone_df$trn_rings = log(abalone_df$rings+1.495)
skewness(abalone_df$trn_rings)
```

```
## [1] 0.0008371152
```

The distribution of Log transformation of the rings

```
library(ggplot2)
ggplot(aes(x = trn_rings), data = abalone_df) +
  geom_histogram(fill="Navy blue") +
  xlab("Number of log(rings)")

## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



```
ggttitle("The Frequency Distribution of Log of Abalone Rings")+
  theme_classic() +
  theme(plot.title = element_text(hjust = 0.5))
```

```
## NULL
```

Handling the outlier data

```
find_outliers = function(x){
  q1 = quantile(x, probs = 0.25)
  q3 = quantile(x, probs = 0.75)

  IQR = q3 - q1

  status = x < q1 - (IQR*1.5) | x > q3 + (IQR*1.5)
}

remove_outliers=function(processed_data,cols = names(processed_data)) {
  for (col in cols) {
    processed_data = processed_data[!find_outliers(processed_data[[col]]), ]
  }
  return(processed_data)
}
```

```

abalone_df = remove_outliers(abalone_df, cols =
                                c("length", "diameter", "height", "weight.whole",
                                  "weight.shucked", "weight.viscera", "weight.shell"))
dim(abalone_df)

```

```
## [1] 4013   10
```

According to the results, current dataset without outliers has 4013 records, thus only 164 records have been removed due to outlier detection.

Convert the categorical variable *sex* into numeric variable

```

# Convert sex to factor
sex_factor <- as.factor(abalone_df$sex)

# Convert factor levels to numeric values
abalone_df[["sex"]] = as.numeric(sex_factor)

```

Data Normalization

```

# Standard Scaling the data
abalone_df = scale(abalone_df, scale = TRUE, center = TRUE)

```

Split the data set in to training and testing set

```

set.seed(2023)

# Remove the 9th column (rings) from the dataset
abalone_df <- subset(abalone_df, select = -rings)

# Split data into training & testing set
sample <- sample(c(TRUE, FALSE), nrow(abalone_df), replace=TRUE, prob=c(0.80,0.20))
train_set <- abalone_df[sample, ]
test_set <- abalone_df[!sample, ]

# Separate predictor and response data from training set
x_train <- train_set[,-9]
y_train <- train_set[,9]

# Separate predictor and response data from testing set
x_test <- test_set[,-9]
y_test <- test_set[,9]

```

Fit the multiple linear regression model with training data including all predictor variables (Transformed response variable - log(rings))

---

```

# -----Multiple Linear Regression -----
# The start time to train the model
start = Sys.time()
# Train Multiple Linear regression model

```

```

ml_model <- lm(trn_rings~., data = as.data.frame(train_set))
# End time to finish the training
end = Sys.time()
# Time taken to train the model
mlr.time = end-start

summary(ml_model)

##
## Call:
## lm(formula = trn_rings ~ ., data = as.data.frame(train_set))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.80733 -0.45401 -0.07893  0.36652  2.81267
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 0.003169  0.011734  0.270   0.787
## sex          0.009402  0.011799  0.797   0.426
## length        0.012161  0.072853  0.167   0.867
## diameter     0.415778  0.073669  5.644 1.81e-08 ***
## height        0.300328  0.029370 10.226 < 2e-16 ***
## weight.whole  1.101572  0.117733  9.357 < 2e-16 ***
## weight.shucked -1.262270  0.060354 -20.914 < 2e-16 ***
## weight.viscera -0.284688  0.047215 -6.030 1.83e-09 ***
## weight.shell    0.313982  0.054937  5.715 1.19e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6678 on 3231 degrees of freedom
## Multiple R-squared:  0.5629, Adjusted R-squared:  0.5619
## F-statistic: 520.2 on 8 and 3231 DF,  p-value: < 2.2e-16

```

The adjusted R squared value of training set is 56.2% with a RMSE of 66.7%. According to the summary results, length and diameter are not significant variables of the model.

Making predictions from the trained multiple linear regression model on the testing set

```

library(ModelMetrics)

##
## Attaching package: 'ModelMetrics'

## The following object is masked from 'package:base':
##
##      kappa

# R squared on train data
mlr.tr.Rsquare = summary(ml_model)$adj.r.squared
# RMSE on train data
mlr.tr.RMSE = 0.6678

```

```

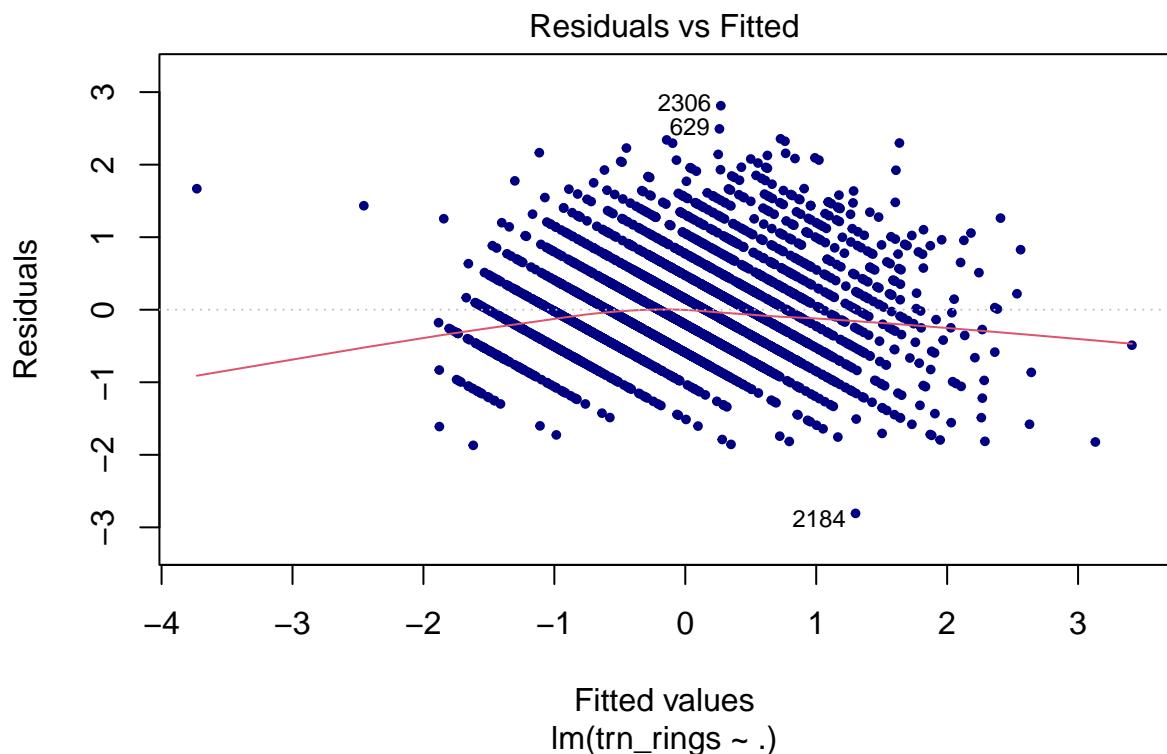
# Predict on test data
y_pred <- predict(ml_model, newdata=as.data.frame(x_test))

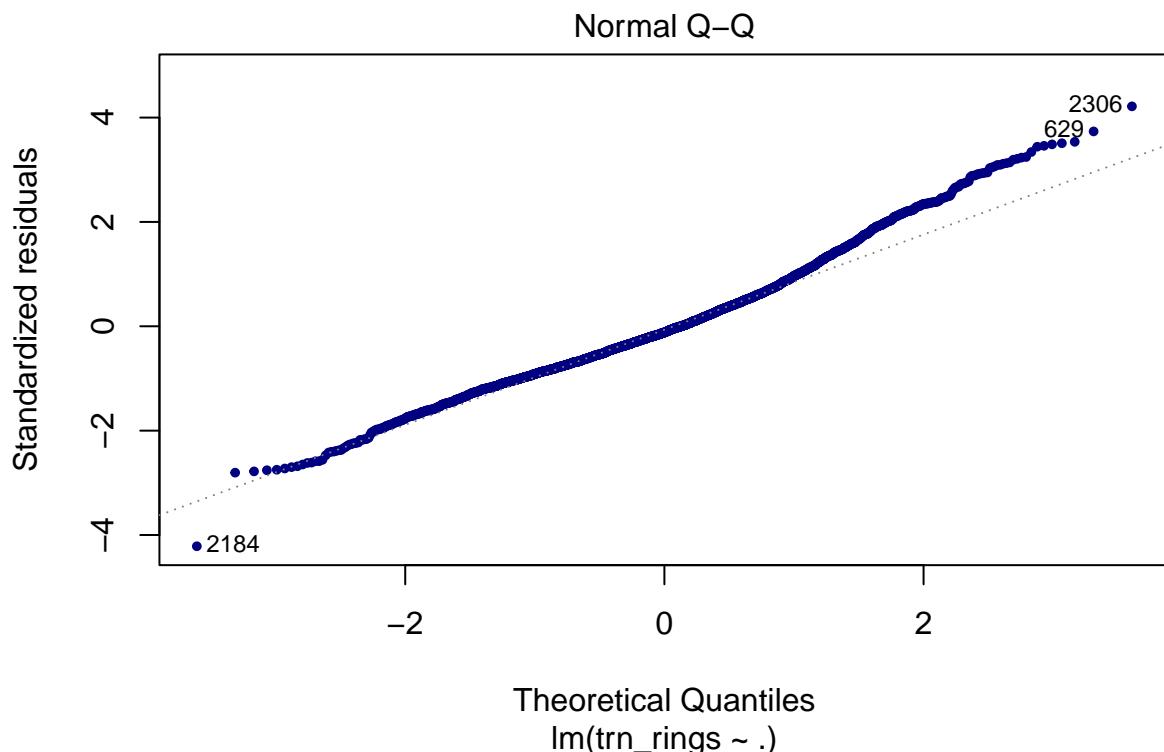
# RMSE and R squared for test set
data.frame(
  mlr.te.RMSE = caret::RMSE(y_pred, y_test),
  mlr.te.Rsquare = caret::R2(y_pred,y_test)
)

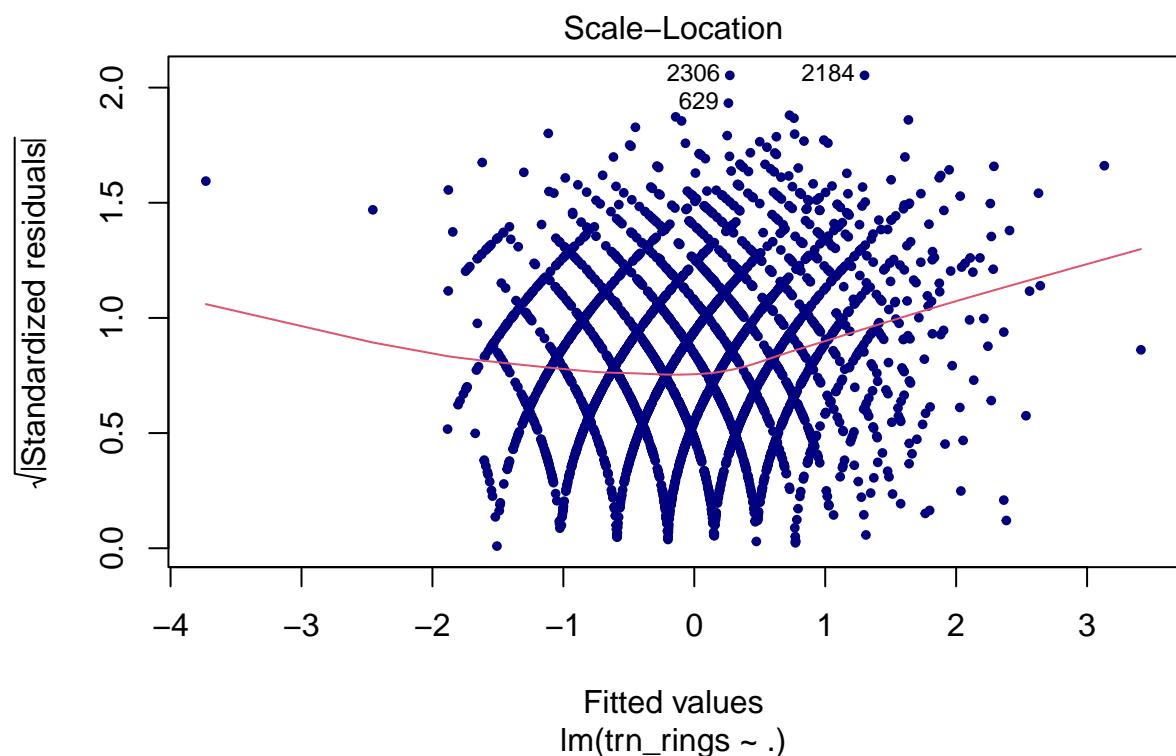
##   mlr.te.RMSE mlr.te.Rsquare
## 1  0.6760124  0.5064648

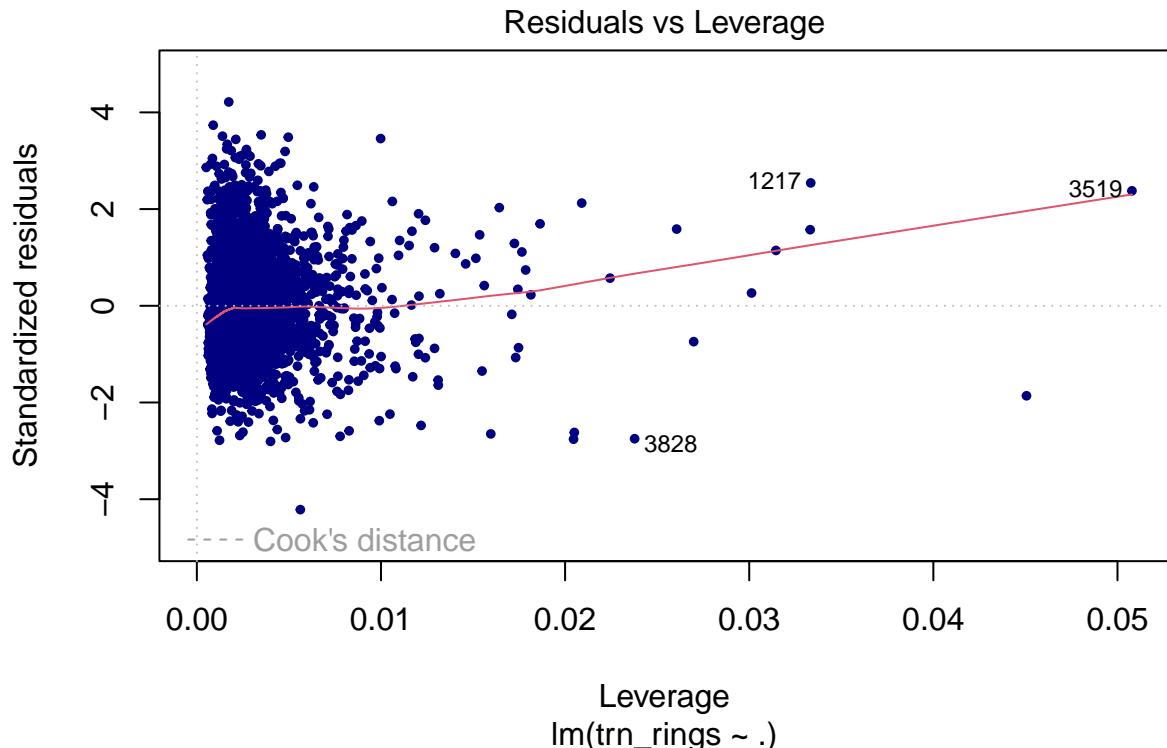
plot(ml_model, col = "navy blue", cex = 0.8, pch=20)

```









We can see that the assumptions of the linear regression is violated and therefore, we first need to remove multicollinearity among the predictor variables, so that the correlation among the predictors will be reduced to a considerable amount to get a better fitted model

To remove the multicollinearity, I would choose principal component regressor, ridge regularization and lasso regularization

PCR - Reducing the dimensions would transform the correlated predictor variables into uncorrelated principal components which would reduce the multicollinearity.

The optimum number of components in the PCR is obtained from hyper parameter tuning

```
# Principal Component Regression with hyper parameter tuning using 10-fold cross validation

library(caret)

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following objects are masked from 'package:ModelMetrics':
##      confusionMatrix, precision, recall, sensitivity, specificity
```

```

# Start time to train the PCR
start = Sys.time()

# define hyperparameters for the number of components
hyperparams <- expand.grid(ncomp = c(1, 2, 3, 4, 5, 6, 7, 8))

# Train PCR model using cross validation
pcr.model <- train(trn_rings~, data = as.data.frame(train_set), method = "pcr",
                     scale = TRUE,
                     trControl = trainControl("cv", number = 10),
                     tuneGrid=hyperparams
)
# End time to finish the training
end = Sys.time()
# Time taken to train the model
pcr.time_taken = end-start

```

```

# Print the best tuning parameter (number of components) that minimize the cross-validation error, RMSE
pcr.model$bestTune

```

```

## ncomp
## 8     8

```

```

# Summarize the final PCR model
summary(pcr.model$finalModel)

```

```

## Data:      X dimension: 3240 8
## Y dimension: 3240 1
## Fit method: svdpc
## Number of components considered: 8
## TRAINING: % variance explained
##           1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X          81.28    93.77   96.16   97.79   98.91   99.74   99.91
## .outcome   37.03    37.05   49.41   49.50   54.50   54.89   54.99
##           8 comps
## X          100.00
## .outcome   56.29

```

According to the results, optimum number of components selected is 8, and 56% of the variability in the log rings can be explained by the predictor variables.

Making predictions on the training and testing data

```

library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
## 
##     filter, lag

```

```

## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union

# Get the R squared and RMSE on training set
train_pred <- pcr.model %>% predict(x_train)
data.frame(
  pcr.tr.RMSE = caret::RMSE(train_pred, y_train),
  pcr.tr.Rsquare = caret::R2(train_pred,y_train)
)

##   pcr.tr.RMSE pcr.tr.Rsquare
## 1  0.6669062      0.5629435

# Make predictions on test data
test_pred <- pcr.model %>% predict(x_test)
data.frame(
  pcr.te.RMSE = caret::RMSE(test_pred, y_test),
  pcr.te.Rsquare = caret::R2(test_pred,y_test)
)

##   pcr.te.RMSE pcr.te.Rsquare
## 1  0.6760124      0.5064648

```

Ridge Regression model

A ridge regression model is fitted with hyperparameter tuning of lambda using 10-fold cross validation

```

library(glmnet)

## Loading required package: Matrix

## Loaded glmnet 4.1-6

# Start time
start = Sys.time()

# define grid of hyper parameters
grid <- expand.grid(alpha = 0, lambda = seq(0, 1, by = 0.01))

# implement 10-fold cross validation
ctrl <- trainControl(method = "cv", number = 10)

# train the ridge regression model using 10-fold cross validation
ridge.model <- train(trn_rings ~ ., data = train_set, method = "glmnet", trControl = ctrl, tuneGrid = grid)

# Optimized alpha
ridge.model$bestTune$alpha

## [1] 0

```

```

# Optimized lambda
ridge.model$bestTune$lambda

## [1] 0.06

# Fit the final model on the training data using optimal lambda
ridge.model <- glmnet(x = as.matrix(x_train), y = as.matrix(y_train), alpha = 0, lambda = ridge.model$bestTune$lambda)

# End time
end = Sys.time()
# Time taken to train the ridge model
ridge.time_taken = end - start

```

Making predictions on the training data and testing data using ridge regression model

```

library(dplyr)

# Make predictions on train_set
train_pred <- ridge.model %>% predict(as.matrix(x_train))
data.frame(
  ridge.tr.RMSE = caret::RMSE(train_pred, y_train),
  ridge.tr.Rsquared = caret::R2(train_pred,y_train)
)

##   ridge.tr.RMSE      s0
## 1     0.687204 0.5395245

# Make predictions on the test data
y_test = as.matrix(y_test)
test_pred <- ridge.model %>% predict(as.matrix(x_test))
data.frame(
  ridge.te.RMSE = caret::RMSE(test_pred, y_test),
  ridge.te.Rsquared = caret::R2(test_pred,y_test)
)

##   ridge.te.RMSE      s0
## 1     0.6887168 0.4870175

```

Lasso Regression

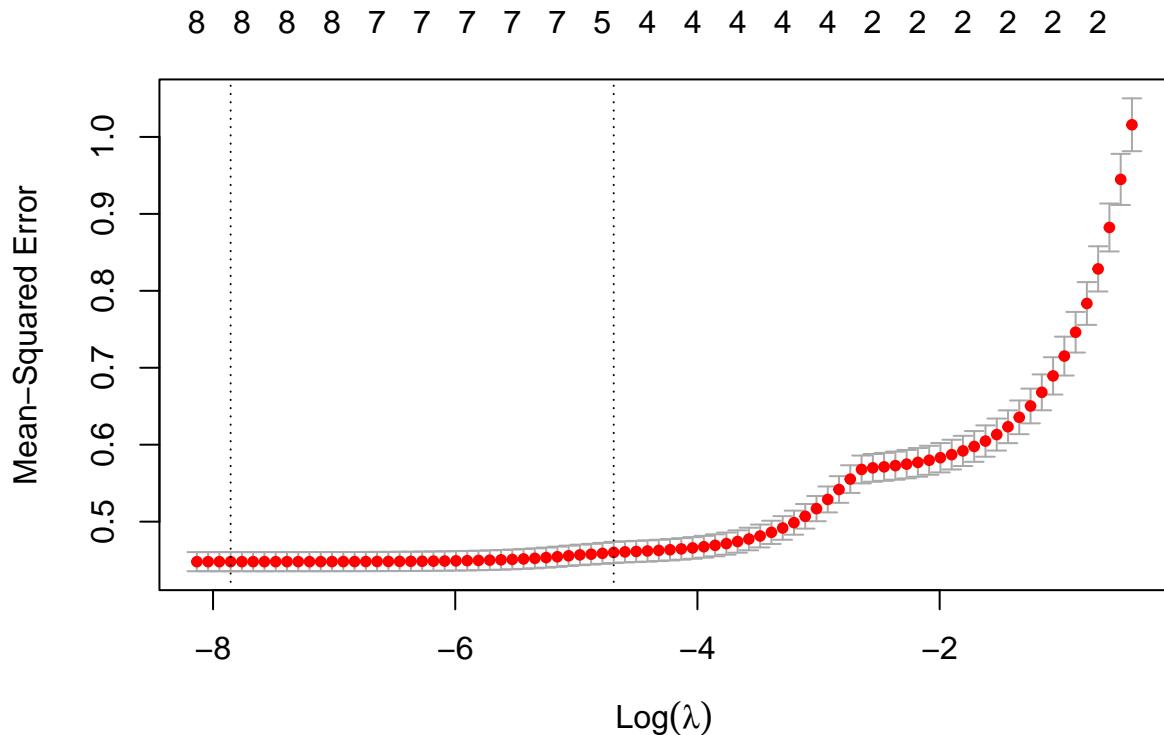
Lasso Regressor is fitted with hyper parameter tuning of lambda using 10-fold cross validation

```

# Start time
start = Sys.time()

# Train the lasso model
cv.lasso <- cv.glmnet(x = x_train, y = as.matrix(y_train), alpha =1, lambda = NULL)
plot(cv.lasso)

```



```

# Optimal lambda
cv.lasso$lambda.min

## [1] 0.0003879473

# Fit the final model on the training data using optimal lambda
lasso.model <- glmnet(x = x_train, y = as.matrix(y_train), alpha = 1, lambda = cv.lasso$lambda.min)

# End time
end = Sys.time()
# Time taken to train the model
lasso.time_taken = end - start

```

The optimal lambda obtained from hyper parameter tuning for lasso regression model is 0.00029

The trained model on optimum lambda is used to make the predictions on training and testing data

```

train_pred <- lasso.model %>% predict(x_train)
data.frame(
  lasso.tr.RMSE = caret::RMSE(train_pred, y_train),
  lasso.tr.Rsquared = caret::R2(train_pred,y_train)
)

```

```

##   lasso.tr.RMSE      s0
## 1     0.6669464 0.5628945

```

```

# Make predictions on the test data

predictions <- lasso.model %>% predict(x_test)
data.frame(
  lasso.te.RMSE = caret::RMSE(predictions, y_test),
  lasso.te.Rsquared = caret::R2(predictions,y_test)
)

##   lasso.te.RMSE      s0
## 1     0.6758612 0.506529

```

### Support Vector Regressor

A support vector regressor is fitted with tuning the hyperparameters (C,Sigma) using 10-fold cross validation

```

#Load Library
library(e1071)

## 
## Attaching package: 'e1071'

## The following objects are masked from 'package:moments':
## 
##      kurtosis, moment, skewness

library(caret)

# Start time
start = Sys.time()
# Define grid for hyper parameter tuning
grid <- expand.grid(sigma = c(0.01, 0.1, 1),
                     C = c(0.1, 1, 10))

# train svm model with cross validation & defined grid of hyper parameters
ctrl <- trainControl(method = "cv", number = 10)
svm.model <- train(trn_rings ~ ., data = train_set, method = "svmRadial", trControl = ctrl, tuneGrid = grid)

# Optimum parameters
best_c <- svm.model$bestTune$C
best_sigma <- svm.model$bestTune$sigma

# train svm model with optimum parameters
svm.model <- train(trn_rings ~ ., data = train_set, kernel = "radial", cost = best_c, epsilon = best_sigma)

# end time
end = Sys.time()
# Time taken to train
svm.time_taken = end - start

# Prediction on training data
train_pred = predict(svm.model, x_train)

```

```

data.frame(
  svm.tr.RMSE = caret::RMSE(train_pred, y_train),
  svm.tr.Rsquare = caret::R2(train_pred, y_train)
)

##   svm.tr.RMSE svm.tr.Rsquare
## 1  0.2998703    0.927161

#Prediction on testing data
test_pred = predict(svm.model, x_test)
data.frame(
  svm.te.RMSE = caret::RMSE(test_pred, y_test),
  svm.te.Rsquare = caret::R2(test_pred, y_test)
)

##   svm.te.RMSE svm.te.Rsquare
## 1  0.630382    0.5711228

# Optimum parameters
best_c

## [1] 1

best_sigma

```

```
## [1] 0.1
```

This is computationally expensive method.

Random Forest Regressor

```

# -----Random Forest Regressor -----
library(caret)

# Start time
start = Sys.time()

train.control <- trainControl(method = "cv", number = 10)
# Train the model
rf.model <- train(trn_rings ~., data = train_set, method = "rf",
                   trControl = train.control,
                   tuneLength=20)

## note: only 7 unique complexity parameters in default grid. Truncating the grid to 7 .

print(rf.model) # Summarize the results

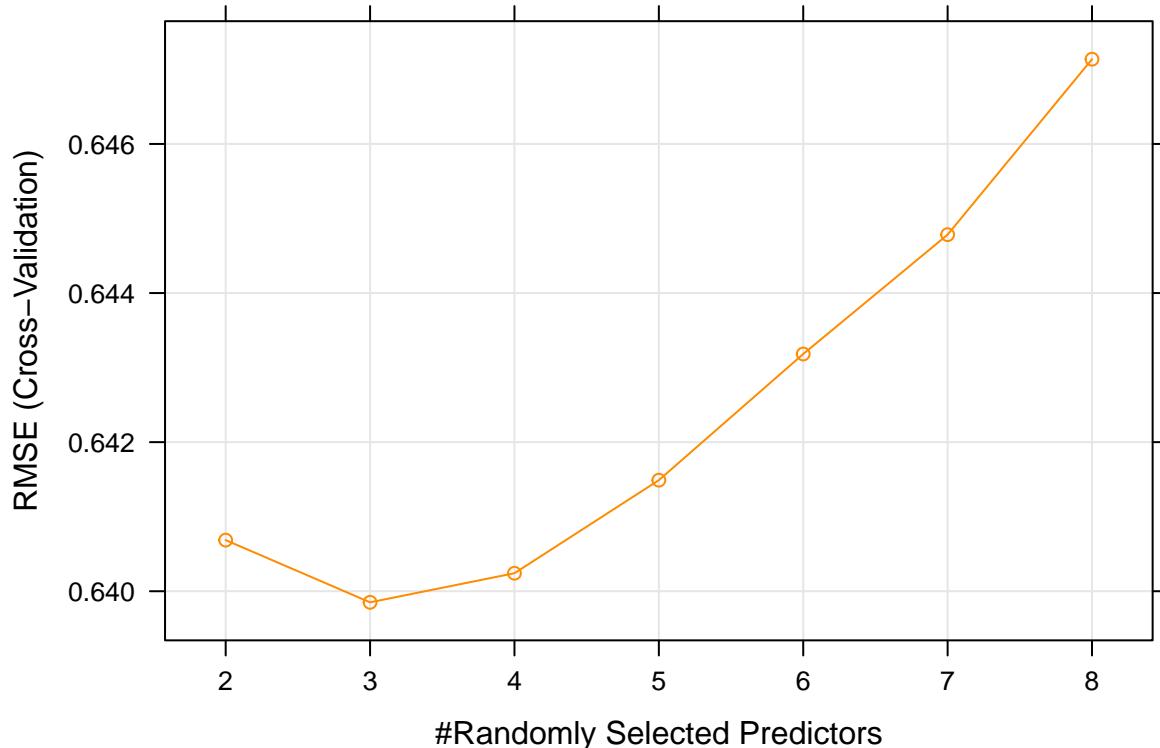
```

```

## Random Forest
##
## 3240 samples
##     8 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 2916, 2916, 2916, 2917, 2916, 2915, ...
## Resampling results across tuning parameters:
##
##   mtry   RMSE      Rsquared    MAE
##   2      0.6406857 0.5963796  0.4913078
##   3      0.6398510 0.5972318  0.4905082
##   4      0.6402411 0.5968272  0.4904695
##   5      0.6414899 0.5953720  0.4919011
##   6      0.6431821 0.5932803  0.4933833
##   7      0.6447835 0.5914076  0.4944942
##   8      0.6471372 0.5887270  0.4964135
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was mtry = 3.

```

```
plot(rf.model, col = "dark orange")
```



```

# end time
end = Sys.time()
# Execution time for RF
rf.time = end-start

# Make predictions on train data
train_pred <- rf.model %>% predict(x_train)

data.frame(
  RMSE = caret::RMSE(train_pred, y_train),
  Rsquare = caret::R2(train_pred, y_train)
)

```

```

##          RMSE    Rsquare
## 1 0.2871294 0.9324538

```

```

# Make predictions on test data
test_pred <- rf.model %>% predict(x_test)

data.frame(
  RMSE = caret::RMSE(test_pred, y_test),
  Rsquare = caret::R2(test_pred, y_test)
)

```

```

##          RMSE    Rsquare
## 1 0.6288769 0.5741445

```

```
varImp(rf.model)
```

```

## rf variable importance
##
##                      Overall
## weight.shell      100.00
## weight.whole      43.92
## height            34.32
## weight.shucked    34.08
## weight.viscera    33.22
## diameter          29.50
## length             13.28
## sex                 0.00

```

```
varImp(rf.model, conditional = TRUE)
```

```

## rf variable importance
##
##                      Overall
## weight.shell      100.00
## weight.whole      43.92
## height            34.32
## weight.shucked    34.08
## weight.viscera    33.22

```

```

## diameter      29.50
## length       13.28
## sex          0.00

```

Decision Tree Regressor

```

# Load libraries
library(rpart)
library(rpart.plot)
library(caret)

# Set up cross-validation
ctrl <- trainControl(method = "cv", number = 10)

# Define hyperparameter grid
grid <- expand.grid(cp = seq(0, 0.1, by = 0.01))

# Start time
start = Sys.time()

# Tune hyperparameters using cross-validation
dt_model <- train(trn_rings ~ .,
                    data = train_set,
                    method = "rpart",
                    trControl = ctrl,
                    tuneGrid = grid)

# Print the best parameters
print(dt_model$bestTune)

##     cp
## 1  0

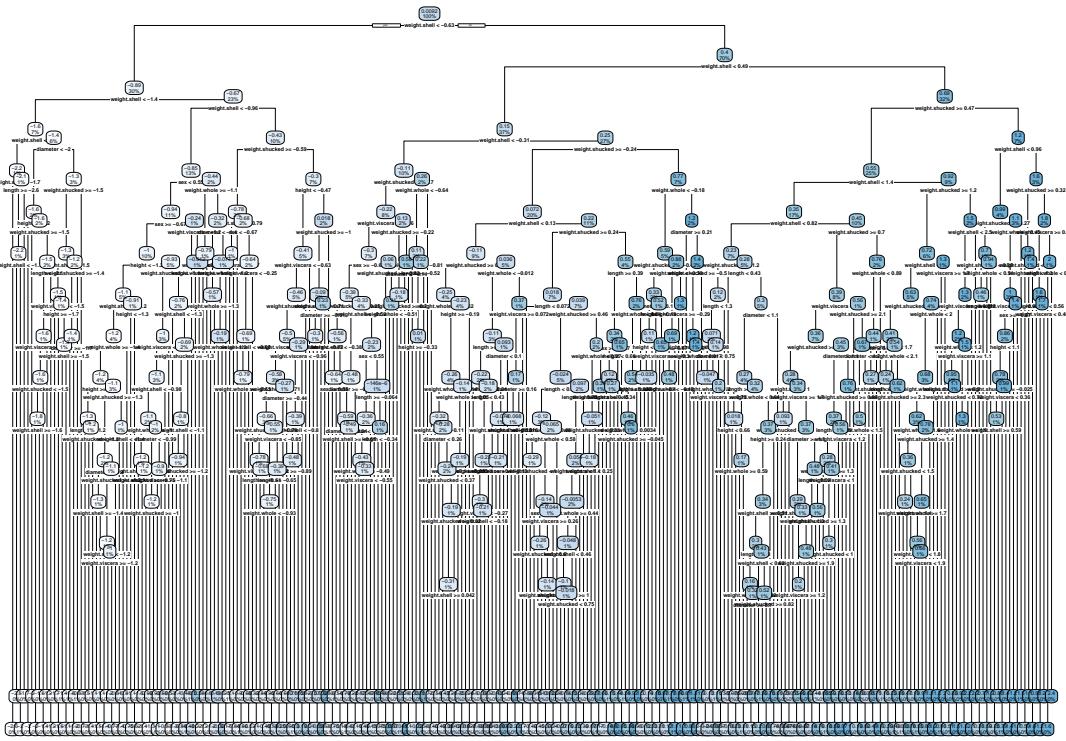
# Train the final model using the best parameters
final_dt_model <- rpart(trn_rings ~ .,
                        data = as.data.frame(train_set),
                        cp = dt_model$bestTune$cp)

# end time
end = Sys.time()
# Execution time for RF
dt.time = end-start

# Plot the tree
rpart.plot(final_dt_model)

## Warning: labs do not fit even at cex 0.15, there may be some overplotting

```



```
# Make predictions on train data
train_pred <- final_dt_model %>% predict(as.data.frame(x_train))
```

```
data.frame(
  RMSE = caret::RMSE(train_pred, y_train),
  Rsquare = caret::R2(train_pred, y_train)
)
```

```
##          RMSE      Rsquare
## 1 0.5086363 0.7457724
```

```
# Make predictions on test data
test_pred <- final_dt_model %>% predict(as.data.frame(x_test))
```

```
data.frame(
  RMSE = caret::RMSE(test_pred, y_test),
  Rsquare = caret::R2(test_pred, y_test)
)
```

```
##          RMSE      Rsquare
## 1 0.714013 0.4849381
```

Neural Network

```

library(nnet)

# Start time
start = Sys.time()

ctrl <- trainControl(method = "cv", number = 10,
                      search = "grid", summaryFunction = defaultSummary)

grid <- expand.grid(size = c(5,6,7,8), decay = seq(0.1, 0.2, 0.3))

model <- train(trn_rings ~ ., data = train_set, method = "nnet", tuneGrid = grid,
               trControl = ctrl)

## # weights:  51
## initial value 3805.133671
## iter  10 value 2631.759919
## iter  20 value 2264.725203
## iter  30 value 2176.536821
## iter  40 value 2163.829767
## iter  50 value 2160.169408
## iter  60 value 2154.428093
## iter  70 value 2152.364967
## iter  80 value 2149.875925
## iter  90 value 2148.286592
## iter 100 value 2147.786012
## final value 2147.786012
## stopped after 100 iterations
## # weights:  61
## initial value 3464.724597
## iter  10 value 2979.710850
## iter  20 value 2335.532572
## iter  30 value 2197.164085
## iter  40 value 2166.622917
## iter  50 value 2157.336502
## iter  60 value 2153.673590
## iter  70 value 2149.097289
## iter  80 value 2147.339871
## iter  90 value 2146.367787
## iter 100 value 2145.724310
## final value 2145.724310
## stopped after 100 iterations
## # weights:  71
## initial value 4310.113161
## iter  10 value 2513.332709
## iter  20 value 2252.636416
## iter  30 value 2171.607746
## iter  40 value 2155.687320
## iter  50 value 2150.495469
## iter  60 value 2147.668441
## iter  70 value 2145.707609
## iter  80 value 2143.354747
## iter  90 value 2141.656528
## iter 100 value 2140.223177

```

```

## final value 2140.223177
## stopped after 100 iterations
## # weights: 81
## initial value 3308.711842
## iter 10 value 2896.905578
## iter 20 value 2285.098042
## iter 30 value 2191.422695
## iter 40 value 2158.069098
## iter 50 value 2151.076690
## iter 60 value 2147.925421
## iter 70 value 2144.799881
## iter 80 value 2143.723679
## iter 90 value 2142.517683
## iter 100 value 2140.777077
## final value 2140.777077
## stopped after 100 iterations
## # weights: 51
## initial value 2990.716101
## iter 10 value 2400.415120
## iter 20 value 2244.773368
## iter 30 value 2196.061155
## iter 40 value 2184.387316
## iter 50 value 2181.298422
## iter 60 value 2176.321703
## iter 70 value 2173.624448
## iter 80 value 2170.693851
## iter 90 value 2169.931266
## iter 100 value 2168.792751
## final value 2168.792751
## stopped after 100 iterations
## # weights: 61
## initial value 3635.535289
## iter 10 value 2999.801145
## iter 20 value 2318.737264
## iter 30 value 2197.402622
## iter 40 value 2175.552441
## iter 50 value 2171.181884
## iter 60 value 2168.883849
## iter 70 value 2166.783699
## iter 80 value 2165.545170
## iter 90 value 2164.892267
## iter 100 value 2164.382683
## final value 2164.382683
## stopped after 100 iterations
## # weights: 71
## initial value 2963.372282
## iter 10 value 2249.303609
## iter 20 value 2203.228797
## iter 30 value 2185.587755
## iter 40 value 2181.833456
## iter 50 value 2180.006334
## iter 60 value 2178.322998
## iter 70 value 2169.011625
## iter 80 value 2164.923395

```

```

## iter  90 value 2163.595308
## iter 100 value 2162.128871
## final  value 2162.128871
## stopped after 100 iterations
## # weights:  81
## initial  value 3656.799595
## iter   10 value 3001.232396
## iter   20 value 2314.686961
## iter   30 value 2211.071341
## iter   40 value 2177.761601
## iter   50 value 2168.685848
## iter   60 value 2164.631469
## iter   70 value 2163.891346
## iter   80 value 2163.178393
## iter   90 value 2162.717325
## iter 100 value 2162.484069
## final  value 2162.484069
## stopped after 100 iterations
## # weights:  51
## initial  value 3431.474955
## iter   10 value 2734.585462
## iter   20 value 2242.861796
## iter   30 value 2177.353953
## iter   40 value 2164.658043
## iter   50 value 2162.126726
## iter   60 value 2160.397571
## iter   70 value 2155.914601
## iter   80 value 2151.540614
## iter   90 value 2149.435081
## iter 100 value 2147.187343
## final  value 2147.187343
## stopped after 100 iterations
## # weights:  61
## initial  value 3710.160471
## iter   10 value 2726.803674
## iter   20 value 2292.280196
## iter   30 value 2177.144869
## iter   40 value 2153.410246
## iter   50 value 2142.718299
## iter   60 value 2140.803394
## iter   70 value 2140.282159
## iter   80 value 2140.143797
## iter   90 value 2140.027170
## iter 100 value 2139.941349
## final  value 2139.941349
## stopped after 100 iterations
## # weights:  71
## initial  value 3603.329272
## iter   10 value 2640.359572
## iter   20 value 2221.960328
## iter   30 value 2162.088087
## iter   40 value 2155.723947
## iter   50 value 2153.628229
## iter   60 value 2150.379222

```

```

## iter 70 value 2145.650198
## iter 80 value 2143.549937
## iter 90 value 2142.113021
## iter 100 value 2141.593176
## final value 2141.593176
## stopped after 100 iterations
## # weights: 81
## initial value 3522.153701
## iter 10 value 2687.998029
## iter 20 value 2199.958911
## iter 30 value 2163.135523
## iter 40 value 2151.183735
## iter 50 value 2146.369399
## iter 60 value 2143.750990
## iter 70 value 2139.552919
## iter 80 value 2138.042710
## iter 90 value 2137.250084
## iter 100 value 2136.737948
## final value 2136.737948
## stopped after 100 iterations
## # weights: 51
## initial value 3841.947782
## iter 10 value 2894.850016
## iter 20 value 2264.576076
## iter 30 value 2188.685803
## iter 40 value 2168.791965
## iter 50 value 2152.193166
## iter 60 value 2149.652616
## iter 70 value 2148.053506
## iter 80 value 2147.608522
## iter 90 value 2147.570505
## iter 100 value 2147.566157
## final value 2147.566157
## stopped after 100 iterations
## # weights: 61
## initial value 3849.756480
## iter 10 value 2725.577292
## iter 20 value 2353.757328
## iter 30 value 2189.327112
## iter 40 value 2165.948979
## iter 50 value 2153.960935
## iter 60 value 2146.743571
## iter 70 value 2144.888443
## iter 80 value 2143.313611
## iter 90 value 2142.316996
## iter 100 value 2142.128374
## final value 2142.128374
## stopped after 100 iterations
## # weights: 71
## initial value 3304.266546
## iter 10 value 2886.872229
## iter 20 value 2416.868528
## iter 30 value 2203.606482
## iter 40 value 2164.087926

```

```

## iter 50 value 2153.730865
## iter 60 value 2150.604160
## iter 70 value 2146.066630
## iter 80 value 2143.394081
## iter 90 value 2141.073959
## iter 100 value 2139.988287
## final value 2139.988287
## stopped after 100 iterations
## # weights: 81
## initial value 3089.856524
## iter 10 value 2595.192790
## iter 20 value 2245.936463
## iter 30 value 2173.072813
## iter 40 value 2155.985071
## iter 50 value 2149.623508
## iter 60 value 2147.131275
## iter 70 value 2145.417649
## iter 80 value 2144.384889
## iter 90 value 2143.392692
## iter 100 value 2141.815206
## final value 2141.815206
## stopped after 100 iterations
## # weights: 51
## initial value 4148.102660
## iter 10 value 2455.288705
## iter 20 value 2256.836291
## iter 30 value 2203.244837
## iter 40 value 2191.313450
## iter 50 value 2185.393336
## iter 60 value 2182.883839
## iter 70 value 2180.225902
## iter 80 value 2176.451362
## iter 90 value 2175.361520
## iter 100 value 2175.272836
## final value 2175.272836
## stopped after 100 iterations
## # weights: 61
## initial value 4312.096244
## iter 10 value 3012.820879
## iter 20 value 2314.827209
## iter 30 value 2224.810333
## iter 40 value 2205.845127
## iter 50 value 2193.326731
## iter 60 value 2182.996080
## iter 70 value 2177.570297
## iter 80 value 2175.306110
## iter 90 value 2174.663668
## iter 100 value 2174.513348
## final value 2174.513348
## stopped after 100 iterations
## # weights: 71
## initial value 4283.923930
## iter 10 value 3008.607405
## iter 20 value 2251.249503

```

```

## iter 30 value 2202.725847
## iter 40 value 2185.313033
## iter 50 value 2177.660398
## iter 60 value 2174.633632
## iter 70 value 2173.342971
## iter 80 value 2172.600897
## iter 90 value 2172.419051
## iter 100 value 2172.362056
## final value 2172.362056
## stopped after 100 iterations
## # weights: 81
## initial value 3287.271815
## iter 10 value 2736.323147
## iter 20 value 2295.066221
## iter 30 value 2219.280637
## iter 40 value 2193.843606
## iter 50 value 2181.963275
## iter 60 value 2178.882053
## iter 70 value 2178.092507
## iter 80 value 2176.062094
## iter 90 value 2175.601587
## iter 100 value 2174.730451
## final value 2174.730451
## stopped after 100 iterations
## # weights: 51
## initial value 3717.055018
## iter 10 value 2810.256968
## iter 20 value 2254.339288
## iter 30 value 2187.186339
## iter 40 value 2168.233773
## iter 50 value 2166.018751
## iter 60 value 2164.149089
## iter 70 value 2162.989276
## iter 80 value 2161.457346
## iter 90 value 2160.943741
## iter 100 value 2159.923670
## final value 2159.923670
## stopped after 100 iterations
## # weights: 61
## initial value 3050.890082
## iter 10 value 2790.699352
## iter 20 value 2221.925355
## iter 30 value 2182.002109
## iter 40 value 2170.539046
## iter 50 value 2166.091471
## iter 60 value 2160.436778
## iter 70 value 2155.834669
## iter 80 value 2155.146695
## iter 90 value 2154.707511
## iter 100 value 2153.955695
## final value 2153.955695
## stopped after 100 iterations
## # weights: 71
## initial value 3643.989111

```

```

## iter 10 value 2913.163047
## iter 20 value 2279.364510
## iter 30 value 2191.353082
## iter 40 value 2167.260081
## iter 50 value 2161.010486
## iter 60 value 2157.607073
## iter 70 value 2155.951729
## iter 80 value 2155.232451
## iter 90 value 2153.925874
## iter 100 value 2153.666568
## final value 2153.666568
## stopped after 100 iterations
## # weights: 81
## initial value 3805.607928
## iter 10 value 3003.568541
## iter 20 value 2350.109285
## iter 30 value 2219.459016
## iter 40 value 2184.635957
## iter 50 value 2166.727739
## iter 60 value 2156.810790
## iter 70 value 2154.419355
## iter 80 value 2153.733241
## iter 90 value 2153.299233
## iter 100 value 2152.754811
## final value 2152.754811
## stopped after 100 iterations
## # weights: 51
## initial value 4102.691894
## iter 10 value 2421.901572
## iter 20 value 2248.847702
## iter 30 value 2211.274097
## iter 40 value 2200.902648
## iter 50 value 2190.179137
## iter 60 value 2188.259358
## iter 70 value 2187.400251
## iter 80 value 2186.829796
## iter 90 value 2186.681529
## iter 100 value 2186.673008
## final value 2186.673008
## stopped after 100 iterations
## # weights: 61
## initial value 3962.162425
## iter 10 value 2527.334425
## iter 20 value 2283.324331
## iter 30 value 2214.667245
## iter 40 value 2201.940896
## iter 50 value 2196.890248
## iter 60 value 2190.516640
## iter 70 value 2186.383523
## iter 80 value 2184.581474
## iter 90 value 2183.145686
## iter 100 value 2182.858658
## final value 2182.858658
## stopped after 100 iterations

```

```

## # weights: 71
## initial value 3426.410654
## iter 10 value 2817.090312
## iter 20 value 2259.295127
## iter 30 value 2211.194468
## iter 40 value 2203.583918
## iter 50 value 2190.241602
## iter 60 value 2183.863475
## iter 70 value 2182.787137
## iter 80 value 2181.899080
## iter 90 value 2181.350314
## iter 100 value 2181.176758
## final value 2181.176758
## stopped after 100 iterations
## # weights: 81
## initial value 3285.540119
## iter 10 value 2992.782424
## iter 20 value 2264.228695
## iter 30 value 2215.850623
## iter 40 value 2202.028123
## iter 50 value 2194.714357
## iter 60 value 2190.267326
## iter 70 value 2186.430674
## iter 80 value 2183.865200
## iter 90 value 2182.067790
## iter 100 value 2181.215102
## final value 2181.215102
## stopped after 100 iterations
## # weights: 51
## initial value 4538.391654
## iter 10 value 2335.589611
## iter 20 value 2225.585142
## iter 30 value 2181.972484
## iter 40 value 2173.956119
## iter 50 value 2171.031106
## iter 60 value 2170.528826
## iter 70 value 2170.414269
## iter 80 value 2170.339561
## iter 90 value 2169.272645
## iter 100 value 2163.920211
## final value 2163.920211
## stopped after 100 iterations
## # weights: 61
## initial value 3825.612789
## iter 10 value 2943.336291
## iter 20 value 2532.186135
## iter 30 value 2248.515855
## iter 40 value 2176.851466
## iter 50 value 2165.495339
## iter 60 value 2161.989423
## iter 70 value 2160.050255
## iter 80 value 2158.522069
## iter 90 value 2157.852887
## iter 100 value 2157.178934

```

```

## final value 2157.178934
## stopped after 100 iterations
## # weights: 71
## initial value 4406.987883
## iter 10 value 2977.945042
## iter 20 value 2279.875387
## iter 30 value 2205.681329
## iter 40 value 2179.627337
## iter 50 value 2172.177468
## iter 60 value 2167.418522
## iter 70 value 2160.717755
## iter 80 value 2158.609732
## iter 90 value 2155.309120
## iter 100 value 2153.273779
## final value 2153.273779
## stopped after 100 iterations
## # weights: 81
## initial value 4376.502895
## iter 10 value 2936.851790
## iter 20 value 2288.163990
## iter 30 value 2194.817712
## iter 40 value 2166.181078
## iter 50 value 2156.619436
## iter 60 value 2154.829062
## iter 70 value 2153.946413
## iter 80 value 2152.856486
## iter 90 value 2151.979850
## iter 100 value 2151.591149
## final value 2151.591149
## stopped after 100 iterations
## # weights: 51
## initial value 3590.750543
## iter 10 value 2949.753032
## iter 20 value 2294.002266
## iter 30 value 2228.761717
## iter 40 value 2185.613349
## iter 50 value 2176.019224
## iter 60 value 2172.643999
## iter 70 value 2172.141053
## iter 80 value 2171.919131
## iter 90 value 2171.729213
## iter 100 value 2171.596990
## final value 2171.596990
## stopped after 100 iterations
## # weights: 61
## initial value 3432.478557
## iter 10 value 2920.378636
## iter 20 value 2235.954884
## iter 30 value 2197.395784
## iter 40 value 2187.749732
## iter 50 value 2177.262188
## iter 60 value 2173.490133
## iter 70 value 2171.977947
## iter 80 value 2171.372589

```

```

## iter  90 value 2168.080794
## iter 100 value 2162.481783
## final  value 2162.481783
## stopped after 100 iterations
## # weights:  71
## initial  value 3226.597697
## iter   10 value 2961.581665
## iter   20 value 2484.302086
## iter   30 value 2243.784520
## iter   40 value 2189.240179
## iter   50 value 2177.586895
## iter   60 value 2174.832722
## iter   70 value 2170.803498
## iter   80 value 2164.664175
## iter   90 value 2160.808670
## iter 100 value 2159.188911
## final  value 2159.188911
## stopped after 100 iterations
## # weights:  81
## initial  value 3041.333247
## iter   10 value 2942.736667
## iter   20 value 2346.715402
## iter   30 value 2221.794668
## iter   40 value 2182.687735
## iter   50 value 2172.646961
## iter   60 value 2168.967487
## iter   70 value 2165.663670
## iter   80 value 2163.560660
## iter   90 value 2162.334713
## iter 100 value 2161.317326
## final  value 2161.317326
## stopped after 100 iterations
## # weights:  51
## initial  value 3424.332712
## iter   10 value 2734.060667
## iter   20 value 2329.223221
## iter   30 value 2179.049300
## iter   40 value 2155.508130
## iter   50 value 2141.828787
## iter   60 value 2138.678563
## iter   70 value 2137.607965
## iter   80 value 2137.196939
## iter   90 value 2137.152727
## iter 100 value 2137.147874
## final  value 2137.147874
## stopped after 100 iterations
## # weights:  61
## initial  value 3735.712625
## iter   10 value 2704.147372
## iter   20 value 2307.030353
## iter   30 value 2192.917118
## iter   40 value 2149.504587
## iter   50 value 2143.203020
## iter   60 value 2140.163138

```

```

## iter 70 value 2137.654560
## iter 80 value 2136.726907
## iter 90 value 2136.532392
## iter 100 value 2136.440983
## final value 2136.440983
## stopped after 100 iterations
## # weights: 71
## initial value 3253.797943
## iter 10 value 2697.199983
## iter 20 value 2198.348649
## iter 30 value 2155.939916
## iter 40 value 2138.895821
## iter 50 value 2134.115618
## iter 60 value 2132.793180
## iter 70 value 2132.403147
## iter 80 value 2132.160779
## iter 90 value 2132.072203
## iter 100 value 2131.978868
## final value 2131.978868
## stopped after 100 iterations
## # weights: 81
## initial value 3640.212593
## iter 10 value 2883.333058
## iter 20 value 2228.818174
## iter 30 value 2174.150837
## iter 40 value 2158.048613
## iter 50 value 2147.143978
## iter 60 value 2137.506588
## iter 70 value 2134.703421
## iter 80 value 2133.792319
## iter 90 value 2133.202437
## iter 100 value 2132.744552
## final value 2132.744552
## stopped after 100 iterations
## # weights: 61
## initial value 4039.775727
## iter 10 value 2602.560550
## iter 20 value 2451.186519
## iter 30 value 2409.666674
## iter 40 value 2399.719828
## iter 50 value 2393.943910
## iter 60 value 2392.577486
## iter 70 value 2392.229019
## iter 80 value 2392.001791
## iter 90 value 2391.907162
## iter 100 value 2391.870078
## final value 2391.870078
## stopped after 100 iterations

best_size = model$bestTune$size
best_decay = model$bestTune$decay

nn.model <- nnet(trn_rings ~ ., data = train_set, size = best_size, decay = best_decay, maxit = 1000)

```

```

## # weights:  61
## initial  value 3576.969387
## iter  10 value 2887.945120
## iter  20 value 2548.289797
## iter  30 value 2437.129019
## iter  40 value 2421.454143
## iter  50 value 2410.727033
## iter  60 value 2403.423830
## iter  70 value 2398.443827
## iter  80 value 2395.352966
## iter  90 value 2393.870168
## iter 100 value 2392.866613
## iter 110 value 2391.932289
## iter 120 value 2391.444617
## iter 130 value 2391.219669
## iter 140 value 2391.173615
## iter 150 value 2391.131595
## final  value 2391.128361
## converged

# end time
end = Sys.time()
# Execution time for RF
nn.time = end-start

# Make predictions on train data
train_pred <- nn.model %>% predict(as.data.frame(x_train))

data.frame(
  RMSE = caret::RMSE(train_pred, y_train),
  Rsquare = caret::R2(train_pred, y_train)
)

##          RMSE    Rsquare
## 1 0.8530333 0.4679529

# Make predictions on test data
test_pred <- nn.model %>% predict(as.data.frame(x_test))

data.frame(
  RMSE = caret::RMSE(test_pred, y_test),
  Rsquare = caret::R2(test_pred, y_test)
)

##          RMSE    Rsquare
## 1 0.8439022 0.407513

```

Bayesian Regression

```
suppressPackageStartupMessages(library(rstanarm))
```

```
# Start time
```

```

start = Sys.time()

model_bayes<- stan_glm(trn_rings~, data=as.data.frame(train_set), family = gaussian(),
prior=normal(0, 5),seed=111)

## 
## SAMPLING FOR MODEL 'continuous' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 0.001 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 10 seconds.
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 1: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 1: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 1: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 1: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 1: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 1: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 1: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 1: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 1: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 1: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.426 seconds (Warm-up)
## Chain 1:           0.672 seconds (Sampling)
## Chain 1:           1.098 seconds (Total)
## Chain 1:
## 
## SAMPLING FOR MODEL 'continuous' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 0 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 2: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 2: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 2: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 2: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 2: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 2: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 2: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 2: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 2: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 2: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 2: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 0.475 seconds (Warm-up)
## Chain 2:           0.712 seconds (Sampling)

```

```

## Chain 2:           1.187 seconds (Total)
## Chain 2:
##
## SAMPLING FOR MODEL 'continuous' NOW (CHAIN 3).
## Chain 3:
## Chain 3: Gradient evaluation took 0 seconds
## Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Chain 3: Adjust your expectations accordingly!
## Chain 3:
## Chain 3:
## Chain 3: Iteration: 1 / 2000 [  0%] (Warmup)
## Chain 3: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 3: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 3: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 3: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 3: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 3: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 3: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 3: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 3: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 3: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 3: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 3:
## Chain 3: Elapsed Time: 0.433 seconds (Warm-up)
## Chain 3:           0.816 seconds (Sampling)
## Chain 3:           1.249 seconds (Total)
## Chain 3:
##
## SAMPLING FOR MODEL 'continuous' NOW (CHAIN 4).
## Chain 4:
## Chain 4: Gradient evaluation took 0 seconds
## Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Chain 4: Adjust your expectations accordingly!
## Chain 4:
## Chain 4:
## Chain 4: Iteration: 1 / 2000 [  0%] (Warmup)
## Chain 4: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 4: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 4: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 4: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 4: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 4: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 4: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 4: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 4: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 4: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 4:
## Chain 4: Elapsed Time: 0.453 seconds (Warm-up)
## Chain 4:           0.751 seconds (Sampling)
## Chain 4:           1.204 seconds (Total)
## Chain 4:

```

```

end = Sys.time()
bayes.time = end-start

summary(model_bayes)

##
## Model Info:
##   function: stan_glm
##   family: gaussian [identity]
##   formula: trn_rings ~ .
##   algorithm: sampling
##   sample: 4000 (posterior sample size)
##   priors: see help('prior_summary')
##   observations: 3240
##   predictors: 9
##
## Estimates:
##           mean    sd   10%   50%   90%
## (Intercept) 0.0    0.0   0.0   0.0   0.0
## sex         0.0    0.0   0.0   0.0   0.0
## length      0.0    0.1  -0.1   0.0   0.1
## diameter    0.4    0.1   0.3   0.4   0.5
## height       0.3    0.0   0.3   0.3   0.3
## weight.whole 1.1    0.1   0.9   1.1   1.3
## weight.shucked -1.3   0.1  -1.3  -1.3  -1.2
## weight.viscera -0.3   0.0  -0.3  -0.3  -0.2
## weight.shell  0.3    0.1   0.2   0.3   0.4
## sigma        0.7    0.0   0.7   0.7   0.7
##
## Fit Diagnostics:
##           mean    sd   10%   50%   90%
## mean_PPD 0.0    0.0   0.0   0.0   0.0
##
## The mean_ppd is the sample average posterior predictive distribution of the outcome variable (for de
##
## MCMC diagnostics
##           mcse Rhat n_eff
## (Intercept) 0.0  1.0  3139
## sex         0.0  1.0  5963
## length      0.0  1.0  2010
## diameter    0.0  1.0  2079
## height       0.0  1.0  4229
## weight.whole 0.0  1.0  1476
## weight.shucked 0.0  1.0  1804
## weight.viscera 0.0  1.0  2790
## weight.shell  0.0  1.0  1850
## sigma        0.0  1.0  3499
## mean_PPD    0.0  1.0  3631
## log-posterior 0.1  1.0  1688
##
## For each parameter, mcse is Monte Carlo standard error, n_eff is a crude measure of effective sample

```

posterior\_predictive checks

```

y_pred = predict(model_bayes,newdata = as.data.frame(x_test))
data.frame(
  RMSE = caret::RMSE(y_pred, y_test),
  Rsquare = caret::R2(y_pred, y_test)
)

##           RMSE   Rsquare
## 1 0.6759975 0.5064838

y_pred_train = predict(model_bayes,newdata = as.data.frame(x_train))
data.frame(
  RMSE = caret::RMSE(y_pred_train, y_train),
  Rsquare = caret::R2(y_pred_train, y_train)
)

##           RMSE   Rsquare
## 1 0.6669065 0.5629431

```

The execution time of each model

```
print("Multiple Linear Regression Model")
```

```
## [1] "Multiple Linear Regression Model"
```

```
mlr.time
```

```
## Time difference of 0.02656293 secs
```

```
print("Principal Component Regressor")
```

```
## [1] "Principal Component Regressor"
```

```
pqr.time_taken
```

```
## Time difference of 0.7460201 secs
```

```
print("Ridge Regressor")
```

```
## [1] "Ridge Regressor"
```

```
ridge.time_taken
```

```
## Time difference of 1.282827 secs
```

```
print("Lasso Regressor")
```

```
## [1] "Lasso Regressor"
```

```

lasso.time_taken

## Time difference of 0.4396169 secs

print("Support Vector Regressor")

## [1] "Support Vector Regressor"

svm.time_taken

## Time difference of 27.3901 mins

print("Random Forest Regressor")

## [1] "Random Forest Regressor"

rf.time

## Time difference of 10.05714 mins

print("Decision Tree Regressor")

## [1] "Decision Tree Regressor"

dt.time

## Time difference of 0.9678252 secs

print("Neural Network")

## [1] "Neural Network"

nn.time

## Time difference of 16.44814 secs

print("Bayesian Regression")

## [1] "Bayesian Regression"

bayes.time

## Time difference of 5.031238 secs

```

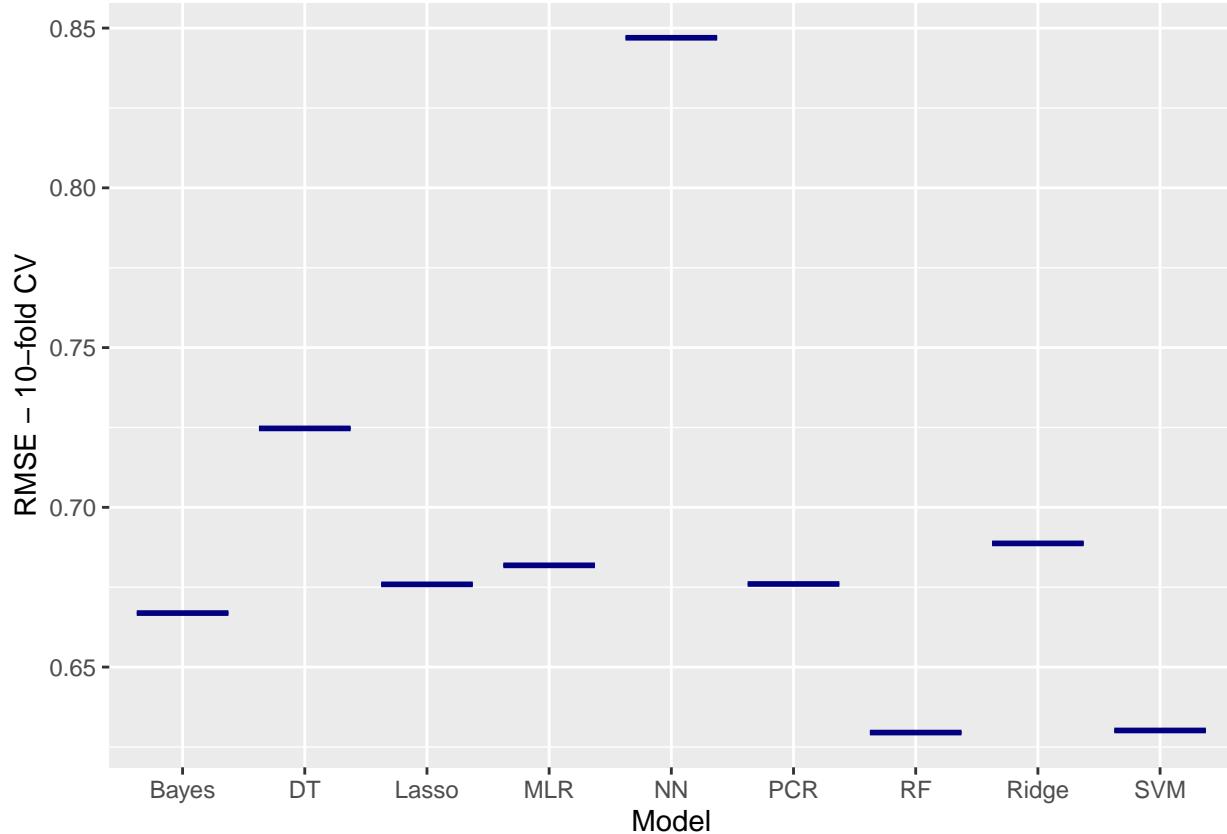
Boxplot representation of RMSE values of each model

```

RMSE_data <- data.frame(Model = c("MLR", "PCR", "Ridge", "Lasso", "SVM", "DT", "RF", "NN", "Bayes"),
                         RMSE = c(0.6818485, 0.6760124, 0.6887168, 0.6758897, 0.6301614, 0.7247068, 0.629501))

ggplot(RMSE_data, aes(x = Model, y = RMSE)) + geom_boxplot(col = "navy blue") + ylab("RMSE - 10-fold CV")

```



The R squared values of all the models are obtained in a boxplot to obtain a clear understanding about the performance of the model.

```

R2_data <- data.frame(Model = c("MLR", "PCR", "Ridge", "Lasso", "SVM", "DT", "RF", "NN", "Bayes"),
                       Rsquare = c(0.4968994, 0.5064648, 0.4870175, 0.5065215, 0.5715283, 0.4373023, 0.629501))

ggplot(R2_data, aes(x = Model, y = Rsquare)) + geom_boxplot(col = "navy blue") + ylab("R-square - 10-fold CV")

```

